# Chapter 3: SQL

Database System Concepts, 5th Ed.

©Silberschatz, Korth and Sudarshan
See **www.db-book.com** for conditions on re-use

Chapter 3: SQL

- Data Definition

- Basic Query Structure

- Set Operations

- Aggregate Functions

- Null Values

- Nested Subqueries

- Complex Queries

- Views

- Modification of the Database

- Joined Relations**

History

- IBM Sequel language developed as part of System R project at the IBM San Jose Research Laboratory

- Renamed Structured Query Language (SQL)

- ANSI and ISO standard SQL:

  - SQL-86

  - SQL-89

  - SQL-92

  - SQL:1999 (language name became Y2K compliant!)

  - SQL:2003

- Commercial systems offer most, if not all, SQL-92 features, plus varying feature sets from later standards and special proprietary features.

 - Not all examples here may work on your particular system.

Data Definition Language

Allows the specification of:

- The schema for each relation, including attribute types.
- Integrity constraints
- Authorization information for each relation.
- Non-standard SQL extensions also allow specification of
  - The set of indices to be maintained for each relations.
 - The physical storage structure of each relation on disk.

Create Table Construct

- An SQL relation is defined using the **create table** command:

  create table r ($A_1 D_1$, $A_2 D_2$, ..., $A_n D_n$,
        (integrity-constraint$_1$),
        ...,
        (integrity-constraint$_k$))

  - r is the name of the relation
  - each A_i is an attribute name in the schema of relation r
 - D_i is the data type of attribute A_i

- Example:

  create table *branch*
      (*branch_name* **char**(15)**,**
      branch_city   **char(**30),
      assets   **integer**)

Domain Types in SQL

- **char(n).** Fixed length character string, with user-specified length n.
- **varchar(n).** Variable length character strings, with user-specified maximum length n.
- **int.** Integer (a finite subset of the integers that is machine-dependent).
- **smallint.** Small integer (a machine-dependent subset of the integer domain type).
- **numeric(p,d).** Fixed point number, with user-specified precision of p digits, with n digits to the right of decimal point.
- **real, double precision.** Floating point and double-precision floating point numbers, with machine-dependent precision.
- **float(n).** Floating point number, with user-specified precision of at least n digits.
- More are covered in Chapter 4.

Integrity Constraints on Tables

- **not null**

- **primary key** $(A_1, ..., A_n)$

Example: Declare *branch_name* as the primary key for *branch*

.

    create table *branch*

          (*branch_name* char(15)**,**
          branch_city   char(30) **not null**,
          assets   integer,
   **primary key** (*branch_name*))

primary key declaration on an attribute automatically ensures **not null**
    in SQL-92 onwards, needs to be explicitly stated in SQL-89

Basic Insertion and Deletion of Tuples

- Newly created table is empty

- Add a new tuple to *account*

    insert into *account*
    values ('A-9732', 'Perryridge', 1200)

    - Insertion fails if any integrity constraint is violated

- Delete *all* tuples from *account*

    delete from *account*

    Note: Will see later how to delete selected tuples

Drop and Alter Table Constructs

- The **drop table** command deletes all information about the dropped relation from the database.

- The **alter table** command is used to add attributes to an existing relation:

   alter table r **add** A D

  where A is the name of the attribute to be added to relation r and D is the domain of A.

   - All tuples in the relation are assigned *null* as the value for the new attribute.

- The **alter table** command can also be used to drop attributes of a relation:

   alter table r **drop** A

  where A is the name of an attribute of relation r

   - Dropping of attributes not supported by many databases

Basic Query Structure

- A typical SQL query has the form:

  select $A_1, A_2, ..., A_n$
  from $r_1, r_2, ..., r_m$
  where P

  - A_i represents an attribute
  - R_i represents a relation
  - P is a predicate.

- This query is equivalent to the relational algebra expression.

$$\prod_{A_1, A_2, \ldots, A_n} (\sigma_P(r_1 \times r_2 \times \ldots \times r_m))$$

- The result of an SQL query is a relation.

The select Clause

- The **select** clause list the attributes desired in the result of a query
  - corresponds to the projection operation of the relational algebra
- Example: find the names of all branches in the *loan* relation:

  select *branch_name*
  from *loan*

- In the relational algebra, the query would be:

$$\prod_{branch\_name} (loan)$$

- NOTE: SQL names are case insensitive (i.e., you may use upper- or lower-case letters.)
  - E.g. *Branch_Name* ≡ *BRANCH_NAME* ≡ *branch_name*
  - Some people use upper case wherever we use bold font.

The select Clause (Cont.)

- SQL allows duplicates in relations as well as in query results.

- To force the elimination of duplicates, insert the keyword **distinct** after select**.**

- Find the names of all branches in the *loan* relations, and remove duplicates

  select distinct *branch_name*
  from *loan*

- The keyword **all** specifies that duplicates not be removed.

 **select all** *branch_name*
  from *loan*

The select Clause (Cont.)

- An asterisk in the select clause denotes "all attributes"

  select *
  from *loan*

- The **select** clause can contain arithmetic expressions involving the operation, +, –, *, and /, and operating on constants or attributes of tuples.

- E.g.:

  select *loan_number, branch_name, amount* * 100
 **from** *loan*

The where Clause

- The **where** clause specifies conditions that the result must satisfy
  - Corresponds to the selection predicate of the relational algebra.
- To find all loan number for loans made at the Perryridge branch with loan amounts greater than $1200.

  select *loan_number*
  from *loan*
  where *branch_name* = 'Perryridge' **and** *amount* > 1200

- Comparison results can be combined using the logical connectives **and, or,** and **not.**

The from Clause

- The **from** clause lists the relations involved in the query
    - Corresponds to the Cartesian product operation of the relational algebra.
- Find the Cartesian product *borrower X loan*

    select *
    from *borrower, loan*

- Find the name, loan number and loan amount of all customers having a loan at the Perryridge branch.

    select *customer_name, borrower.loan_number, amount*
        from *borrower, loan*
        where *borrower.loan_number = loan.loan_number* **and**
   *branch_name* = 'Perryridge'

The Rename Operation

- SQL allows renaming relations and attributes using the **as** clause:

   old-name **as** *new-name*

- E.g. Find the name, loan number and loan amount of all customers; rename the column name *loan_number* as *loan_id.*

   select *customer_name, borrower.loan_number* **as** *loan_id, amount*
   from *borrower, loan*
   where *borrower.loan_number = loan.loan_number*

Tuple Variables

- Tuple variables are defined in the **from** clause via the use of the **as** clause.

- Find the customer names and their loan numbers and amount for all customers having a loan at some branch.

$$\textbf{select } \textit{customer\_name, T.loan\_number, S.amount}$$
$$\textbf{from } \textit{borrower } \textbf{as } \textit{T, loan } \textbf{as } \textit{S}$$
$$\textbf{where } \textit{T.loan\_number = S.loan\_number}$$

- Find the names of all branches that have greater assets than some branch located in Brooklyn.

$$\textbf{select distinct } \textit{T.branch\_name}$$
$$\textbf{from } \textit{branch } \textbf{as } \textit{T, branch } \textbf{as } \textit{S}$$
$$\textbf{where } \textit{T.assets > S.assets } \textbf{and } \textit{S.branch\_city = } \textit{'Brooklyn'}$$

- Keyword **as** is optional and may be omitted

$$\textit{borrower } \textbf{as } \textit{T} \equiv \textit{borrower T}$$

 - Some database such as Oracle *require* **as** to be omitted

String Operations

- SQL includes a string-matching operator for comparisons on character strings. The operator "like" uses patterns that are described using two special characters:

  - percent (%). The % character matches any substring.

  - underscore (_). The _ character matches any character.

- Find the names of all customers whose street includes the substring "Main".

  select *customer_name*
  from *customer*
  where *customer_street* **like '**% Main%'

- Match the name "Main%"

  like 'Main\%' **escape '**\'

- SQL supports a variety of string operations such as

  - concatenation (using "||")

  - converting from upper to lower case (and vice versa)

 - finding string length, extracting substrings, etc.

Ordering the Display of Tuples

- List in alphabetic order the names of all customers having a loan in Perryridge branch

    select distinct *customer_name*
    from *borrower, loan*
    where *borrower loan_number = loan.loan_number* **and**
    branch_name = 'Perryridge'
   **order by** *customer_name*

- We may specify **desc** for descending order or **asc** for ascending order, for each attribute; ascending order is the default.

    - Example: **order by** *customer_name* **desc**

Duplicates

- In relations with duplicates, SQL can define how many copies of tuples appear in the result.

- **Multiset** versions of some of the relational algebra operators – given multiset relations r_1 and r_2:

  1. $\sigma_\theta(r_1)$: If there are c_1 copies of tuple t_1 in r_1, and t_1 satisfies selections $\sigma_{\theta,}$ then there are c_1 copies of t_1 in $\sigma_\theta(r_1)$.

  2. $\Pi_A(r)$: For each copy of tuple t_1 in r_1, there is a copy of tuple $\Pi_A(t_1)$ in $\Pi_A(r_1)$ where $\Pi_A(t_1)$ denotes the projection of the single tuple t_1.

  3. $r_1 \times r_2$: If there are c_1 copies of tuple t_1 in r_1 and c_2 copies of tuple t_2 in r_2, there are $c_1 \times c_2$ copies of the tuple $t_1. t_2$ in $r_1 \times r_2$

Duplicates (Cont.)

- Example: Suppose multiset relations r_1 (A, B) and r_2 (C) are as follows:

  $r_1 = \{(1, a)\ (2,a)\} \quad r_2 = \{(2), (3), (3)\}$

- Then $\Pi_B(r_1)$ would be $\{(a), (a)\}$, while $\Pi_B(r_1) \times r_2$ would be

  $\{(a,2), (a,2), (a,3), (a,3), (a,3), (a,3)\}$

- SQL duplicate semantics:

  select $A_1, A_2, ..., A_n$
  from $r_1, r_2, ..., r_m$
  where P

  is equivalent to the *multiset* version of the expression:

$$\Pi_{A_1, A_2, \ldots, A_n}(\sigma_P(r_1 \times r_2 \times \cdots \times r_m))$$

Set Operations

- The set operations **union**, **intersect**, and **except** operate on relations and correspond to the relational algebra operations $\cup, \cap, -$.

- Each of the above operations automatically eliminates duplicates; to retain all duplicates use the corresponding multiset versions **union all**, **intersect all** and **except all**.

  Suppose a tuple occurs m times in r and n times in s, then, it occurs:

  - $m + n$ times in r **union all** s

  - $\min(m,n)$ times in r **intersect all** s

  - $\max(0, m - n)$ times in r **except all** s

Set Operations

- Find all customers who have a loan, an account, or both:

    (**select** *customer_name* **from** *depositor*)
    union
    (**select** *customer_name* **from** *borrower)*

- Find all customers who have both a loan and an account.

    (**select** *customer_name* **from** *depositor*)
    intersect
    (**select** *customer_name* **from** *borrower)*

- Find all customers who have an account but no loan.

    (**select** *customer_name* **from** *depositor*)
    except
    (**select** *customer_name* **from** *borrower)*

Aggregate Functions

- These functions operate on the multiset of values of a column of a relation, and return a value

   avg: average value
   min: minimum value
   max: maximum value
   sum: sum of values
   count: number of values

Aggregate Functions (Cont.)

- Find the average account balance at the Perryridge branch.

  select avg *(balance)*
       from *account*
       where *branch_name* = 'Perryridge'

- Find the number of tuples in the *customer* relation.

  select count (*)
       from *customer*

- Find the number of depositors in the bank.

  select count (distinct *customer_name)*
   **from** *depositor*

Aggregate Functions – Group By

- Find the number of depositors for each branch.

    select *branch_name,* **count (distinct** *customer_name)*
        from *depositor, account*
        where *depositor.account_number = account.account_number*
        group by *branch_name*

    Note: Attributes in **select** clause outside of aggregate functions must
   appear in **group by** list

Aggregate Functions – Having Clause

- Find the names of all branches where the average account balance is more than $1,200.

> **select** *branch_name,* **avg** (*balance*)
>   **from** *account*
>   **group by** *branch_name*
>   **having avg** (*balance*) > 1200

Note: predicates in the **having** clause are applied after the formation of groups whereas predicates in the **where** clause are applied before forming groups

Nested Subqueries

- SQL provides a mechanism for the nesting of subqueries.

- A **subquery** is a **select-from-where** expression that is nested within another query.

- A common use of subqueries is to perform tests for set membership, set comparisons, and set cardinality.

"In" Construct

- Find all customers who have both an account and a loan at the bank.

> **select distinct** *customer_name*
>   **from** *borrower*
>   **where** *customer_name* **in** (**select** *customer_name*
>   **from** *depositor*)

- Find all customers who have a loan at the bank but do not have an account at the bank

> **select distinct** *customer_name*
>   **from** *borrower*
>   **where** *customer_name* **not in** (**select** *customer_name*
>   **from** *depositor*)

Example Query

- Find all customers who have both an account and a loan at the Perryridge branch

  select distinct *customer_name*
      from *borrower, loan*
      where *borrower.loan_number = loan.loan_number* **and**
         branch_name = 'Perryridge' **and**
       (*branch_name, customer_name*) **in**
       (select *branch_name, customer_name*
        from *depositor, account*
        where *depositor.account_number =*
             account.account_number)

- Note: Above query can be written in a much simpler manner. The formulation above is simply to illustrate SQL features.

# "Some" Construct

- Find all branches that have greater assets than some branch located in Brooklyn.

    select distinct *T.branch_name*
      from *branch* **as** *T, branch* **as** *S*
      where *T.assets* > *S.assets* **and**
         S.branch_city = 'Brooklyn'

- Same query using > **some** clause

    select *branch_name*
      from *branch*
      where *assets* > **some**
        (select *assets*
         from *branch*
   **where** *branch_city =* 'Brooklyn')

"All" Construct

- Find the names of all branches that have greater assets than all branches located in Brooklyn.

$$\textbf{select}\ branch\_name$$
$$\qquad \textbf{from}\ branch$$
$$\qquad \textbf{where}\ assets > \textbf{all}$$
$$\qquad\qquad \textbf{(select}\ assets$$
$$\qquad\qquad \textbf{from}\ branch$$
$$\qquad\qquad \textbf{where}\ branch\_city = \text{'Brooklyn')}$$

"Exists" Construct

- Find all customers who have an account at all branches located in Brooklyn.

> **select distinct** *S.customer_name*
>   **from** *depositor* **as** *S*
>   **where not exists** (
>   (**select** *branch_name*
>   **from** *branch*
>   **where** *branch_city* = 'Brooklyn')
>   **except**
>   (**select** *R.branch_name*
>   **from** *depositor* **as** *T, account* **as** *R*
>   **where** *T.account_number* = *R.account_number* **and**
>   *S.customer_name* = *T.customer_name*))

- Note that $X - Y = \emptyset \iff X \subseteq Y$

- *Note:* Cannot write this query using = **all** and its variants

Absence of Duplicate Tuples

- The **unique** construct tests whether a subquery has any duplicate tuples in its result.

- Find all customers who have at most one account at the Perryridge branch.

  select *T.customer_name*
  from *depositor* **as** *T*
  where unique (

     select *R.customer_name*
     from *account, depositor* **as** *R*
     where *T.customer_name* = *R.customer_name* **and**
        R.account_number = *account.account_number* **and**
   *account.branch_name* = 'Perryridge')

Example Query

- Find all customers who have at least two accounts at the Perryridge branch.

> **select distinct** *T.customer_name*
> **from** *depositor* **as** *T*
> **where not unique** (
>   **select** *R.customer_name*
>   **from** *account, depositor* **as** *R*
>   **where** *T.customer_name* = *R.customer_name* **and**
>   *R.account_number* = *account.account_number* **and**
>   *account.branch_name* = 'Perryridge')

- Variable from outer level is known as a **correlation variable**

Modification of the Database – Deletion

- Delete all account tuples at the Perryridge branch

  delete from *account*
  where *branch_name* = 'Perryridge'

- Delete all accounts at every branch located in the city 'Needham'.

  delete from *account*
  where *branch_name* **in** (**select** *branch_name*
  　　　　　from *branch*
 　　　　　**where** *branch_city* = 'Needham')

Example Query

- Delete the record of all accounts with balances below the average at the bank.

> **delete from** *account*
>   **where** *balance* < (**select avg** (*balance*)
>   **from** *account*)

- Problem: as we delete tuples from deposit, the average balance changes
- Solution used in SQL:
    1. First, compute **avg** balance and find all tuples to delete
    2. Next, delete all tuples found above (without recomputing **avg** or retesting the tuples)

Modification of the Database – Insertion

- Add a new tuple to *account*

    insert into *account*
         values ('A-9732', 'Perryridge', 1200)

   or equivalently

     insert into *account* (*branch_name, balance, account_number*)
     values ('Perryridge', 1200, 'A-9732')

- Add a new tuple to *account* with *balance* set to null

    insert into *account*
   **values** ('A-777','Perryridge', *null*)

Modification of the Database – Insertion

- Provide as a gift for all loan customers of the Perryridge branch, a $200 savings account. Let the loan number serve as the account number for the new savings account

  insert into *account*
  select *loan_number, branch_name,* 200
  from *loan*
  where *branch_name* = 'Perryridge'
    insert into *depositor*
  select *customer_name, loan_number*
  from *loan, borrower*
  where branch_name = 'Perryridge'
          and *loan.account_number = borrower.account_number*

- The **select from where** statement is evaluated fully before any of its results are inserted into the relation

  - Motivation: **insert into** *table*1 **select * from** *table*1

Modification of the Database – Updates

- Increase all accounts with balances over $10,000 by 6%, all other accounts receive 5%.

    - Write two **update** statements:

        update *account*
        set *balance = balance* * 1.06
        where *balance* > 10000

        update *account*
        set *balance = balance* * 1.05
        where *balance* ≤ 10000

    - The order is important

   - Can be done better using the **case** statement (next slide)

Case Statement for Conditional Updates

- Same query as before: Increase all accounts with balances over $10,000 by 6%, all other accounts receive 5%.

update *account*
set *balance* = **case**
      when *balance* $<=$ 10000 **then** *balance* $*1.05$
      else *balance* $* 1.06$
   **end**

More Features

Joined Relations**

- **Join operations** take two relations and return as a result another relation.
- These additional operations are typically used as subquery expressions in the **from** clause
- **Join condition** – defines which tuples in the two relations match, and what attributes are present in the result of the join.
- **Join type** – defines how tuples in each relation that do not match any tuple in the other relation (based on the join condition) are treated.

Join types
inner join
left outer join
right outer join
full outer join

Join Conditions
natural
on < predicate >
using $(A_1, A_1, ..., A_n)$

Joined Relations – Datasets for Examples

- Relation *loan*

- Relation *borrower*

loan_number	branch_name	amount
L-170	Downtown	3000
L-230	Redwood	4000
L-260	Perryridge	1700

loan

customer_name	loan_number
Jones	L-170
Smith	L-230
Hayes	L-155

borrower

- Note: borrower information missing for L-260 and loan information missing for L-155

Joined Relations – Examples

- *loan* **inner join** *borrower* **on**
  loan.loan_number = borrower.loan_number

loan_number	branch_name	amount	customer_name	loan_number
L-170	Downtown	3000	Jones	L-170
L-230	Redwood	4000	Smith	L-230

- *loan* **left outer join** *borrower* **on**
  loan.loan_number = borrower.loan_number

loan_number	branch_name	amount	customer_name	loan_number
L-170	Downtown	3000	Jones	L-170
L-230	Redwood	4000	Smith	L-230
L-260	Perryridge	1700	*null*	*null*

Joined Relations – Examples

- *loan* **natural inner join** *borrower*

loan_number	branch_name	amount	customer_name
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith

- *loan* **natural right outer join** *borrower*

loan_number	branch_name	amount	customer_name
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith
L-155	null	null	Hayes

- Find all customers who have either an account or a loan (but not both) at the bank.

select *customer_name*
    from (*depositor* **natural full outer join** *borrower*)
    where *account_number* **is null or** *loan_number* **is null**

Joined Relations – Examples

- Natural join can get into trouble if two relations have an attribute with same name that should not affect the join condition
  - e.g. an attribute such as *remarks* may be present in many tables
- *Solution:*
  - *loan* **full outer join** *borrower* **using** (*loan_number*)

| loan_number	branch_name	amount	customer_name
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith
L-260	Perryridge	1700	null
L-155	null	null	Hayes

Derived Relations

- SQL allows a subquery expression to be used in the **from** clause

- Find the average account balance of those branches where the average account balance is greater than $1200.

    select *branch_name, avg_balance*
    from (**select** *branch_name,* **avg** (*balance*)
                    from *account*
                    group by *branch_name*)
            as *branch_avg* (*branch_name, avg_balance*)
   **where** *avg_balance* > 1200

    Note that we do not need to use the **having** clause, since we compute the temporary (view) relation *branch_avg* in the **from** clause, and the attributes of *branch_avg* can be used directly in the **where** clause.

View Definition

- A relation that is not of the conceptual model but is made visible to a user as a "virtual relation" is called a **view**.

- A view is defined using the **create view** statement which has the form

  create view *v* **as** < query expression >

  where <query expression> is any legal SQL expression. The view name is represented by *v*.

- Once a view is defined, the view name can be used to refer to the virtual relation that the view generates.

Example Queries

- A view consisting of branches and their customers

  create view *all_customer* **as**
      (**select** *branch_name, customer_name*
       from *depositor, account*
      where *depositor.account_number =*
          account.account_number)
      union
      (**select** *branch_name, customer_name*
      from *borrower, loan*
      where *borrower.loan_number = loan.loan_number*)

- Find all customers of the Perryridge branch

  select *customer_name*
      from *all_customer*
   **where** *branch_name* = 'Perryridge'

Uses of Views

- Hiding some information from some users

    - Consider a user who needs to know a customer's name, loan number and branch name, but has no need to see the loan amount.

    - Define a view

            (**create view** *cust_loan_data* **as**
              select *customer_name, borrower.loan_number, branch_name*
               from *borrower, loan*
              where *borrower.loan_number = loan.loan_number*)

    - Grant the user permission to read *cust_loan_data,* but not *borrower or loan*

- Predefined queries to make writing of other queries easier

    - Common example: Aggregate queries used for statistical analysis of data

Processing of Views

- When a view is created
  - the query expression is stored in the database along with the view name
  - the expression is substituted into any query using the view
- Views definitions containing views
  - One view may be used in the expression defining another view
  - A view relation v_1 is said to *depend directly* on a view relation v_2 if v_2 is used in the expression defining v_1
  - A view relation v_1 is said to *depend on* view relation v_2 if either v_1 depends directly to v_2 or there is a path of dependencies from v_1 to v_2
  - A view relation v is said to be *recursive* if it depends on itself.

View Expansion

- A way to define the meaning of views defined in terms of other views.

- Let view v_1 be defined by an expression e_1 that may itself contain uses of view relations.

- View expansion of an expression repeats the following replacement step:

  repeat
        Find any view relation v_i in e_1
        Replace the view relation v_i by the expression defining v_i
    until no more view relations are present in e_1

- As long as the view definitions are not recursive, this loop will terminate

With Clause

- The **with** clause provides a way of defining a temporary view whose definition is available only to the query in which the **with** clause occurs.

- Find all accounts with the maximum balance

  with *max_balance* (*value*) **as**
     select max (*balance*)
     from *account*
  select *account_number*
  from *account, max_balance*
  where *account.balance = max_balance.value*

Complex Queries using With Clause

- Find all branches where the total account deposit is greater than the average of the total account deposits at all branches.

  with *branch_total* (*branch_name*, *value*) **as**
     select *branch_name*, **sum** (*balance*)
     from *account*
     group by *branch_name*
  with *branch_total_avg* (*value*) **as**
     select avg (*value*)
     from *branch_total*
  select *branch_name*
  from *branch_total*, *branch_total_avg*
  where *branch_total.value* >= *branch_total_avg.value*

- Note: the exact syntax supported by your database may vary slightly.

  - *E.g.* Oracle syntax is of the form
    with *branch_total* **as** (**select** ..),
       branch_total_avg **as** (**select** ..)
    select …

Update of a View

- Create a view of all loan data in the *loan* relation, hiding the *amount* attribute

    create view *loan_branch* **as**
    select *loan_number, branch_name*
   **from** *loan*

- Add a new tuple to *loan_branch*

    insert into *loan_branch*
    values ('L-37', 'Perryridge')

    This insertion must be represented by the insertion of the tuple

    ('L-37', 'Perryridge', *null*)

    into the *loan* relation

Updates Through Views (Cont.)

- Some updates through views are impossible to translate into updates on the database relations

  - **create view** *v* **as**
    - **select** *loan_number, branch_name, amount*
      - **from** *loan*
      - **where** *branch_name* = 'Perryridge'

    insert into *v* **values** ('L-99','Downtown', '23')

- Others cannot be translated uniquely

  - **insert into** *all_customer* **values** ('Perryridge', 'John')
    - 4 Have to choose loan or account, and create a new loan/account number!

- Most SQL implementations allow updates only on simple views (without aggregates) defined on a single relation

Null Values

- It is possible for tuples to have a null value, denoted by *null*, for some of their attributes

- *null* signifies an unknown value or that a value does not exist.

- The predicate **is null** can be used to check for null values.

  - Example: Find all loan number which appear in the *loan* relation with null values for *amount*.

 **select** *loan_number*
  from *loan*
  where *amount* **is null**

- The result of any arithmetic expression involving *null* is *null*

  - Example: $5 + null$ returns null

- However, aggregate functions simply ignore nulls

 - More on next slide

Null Values and Three Valued Logic

- Any comparison with *null* returns *unknown*

  - Example*: 5 < null   or   null <> null   or   null = null*

- Three-valued logic using the truth value *unknown*:

  - OR: (*unknown* **or** *true*)   = *true*,
    (*unknown* **or** *false*) = *unknown*
    (*unknown* **or** *unknown*) = *unknown*

  - AND: *(true* **and** *unknown)* = *unknown,*
    (false **and** *unknown) = false,*
    (unknown **and** *unknown) = unknown*

  - NOT*: (***not** unknown) = unknown*

  - "*P* **is unknown**" evaluates to true if predicate *P* evaluates to *unknown*

- Result of **where** clause predicate is treated as *false* if it evaluates to *unknown*

Null Values and Aggregates

- Total all loan amounts

  select sum (*amount*)
  from *loan*

  - Above statement ignores null amounts

  - Result is *null* if there is no non-null amount

- All aggregate operations except **count(*)** ignore tuples with null values on the aggregated attributes.

End of Chapter 3

The where Clause (Cont.)

- SQL includes a **between** comparison operator

- Example: Find the loan number of those loans with loan amounts between $90,000 and $100,000 (that is, ≥ $90,000 and ≤ $100,000)

  select *loan_number*
     from *loan*
   **where** *amount* **between** 90000 **and** 100000

Figure 3.1: Database Schema

branch (*branch_name, branch_city, assets*)

customer (*customer_name, customer_street, customer_city*)

loan (*loan_number, branch_name, amount*)

borrower (*customer_name, loan_number*)

account (*account_number, branch_name, balance*)

depositor (*customer_name, account_number*)

Definition of Some Clause

$(5 = \textbf{some}\ \boxed{\begin{array}{c} 0 \\ \hline 5 \end{array}}\ ) = \text{true}$

$(5 \neq \textbf{some}\ \boxed{\begin{array}{c} 0 \\ \hline 5 \end{array}}\ ) = \text{true (since } 0 \neq 5)$

- $(= \textbf{some}) \equiv \textbf{in}$
- However, $(\neq \textbf{some})$ is not equivalent to **not in**

Definition of all Clause

$(5 < \mathbf{all}\ \boxed{\begin{matrix} 0 \\ 5 \\ 6 \end{matrix}}\ ) = \text{false}$

$(5 < \mathbf{all}\ \boxed{\begin{matrix} 6 \\ 10 \end{matrix}}\ ) = \text{true}$

$(5 = \mathbf{all}\ \boxed{\begin{matrix} 4 \\ 5 \end{matrix}}\ ) = \text{false}$

$(5 \neq \mathbf{all}\ \boxed{\begin{matrix} 4 \\ 6 \end{matrix}}\ ) = \text{true}$ (since $5 \neq 4$ and $5 \neq 6$)

- $(\neq \mathbf{all}) \equiv \mathbf{not\ in}$
- However, $(= \mathbf{all})$ is not equivalent to $\mathbf{in}$

Test for Empty Relations

- The **exists** construct returns the value **true** if the argument subquery is nonempty.

- **exists** $r \Leftrightarrow r \neq \emptyset$

- **not exists** $r \Leftrightarrow r = \emptyset$

Figure 3.3: Tuples inserted into *loan* and *borrower*

loan_number	branch_name	amount
L-11	Round Hill	900
L-14	Downtown	1500
L-15	Perryridge	1500
L-16	Perryridge	1300
L-17	Downtown	1000
L-23	Redwood	2000
L-93	Mianus	500
null	null	1900

loan

customer_name	loan_number
Adams	L-16
Curry	L-93
Hayes	L-15
Jackson	L-14
Jones	L-17
Smith	L-11
Smith	L-23
Williams	L-17
Johnson	null

borrower

Figure 3.4:
The *loan* and *borrower* relations

loan_number	branch_name	amount	customer_name	loan_number
L-170	Downtown	3000	Jones	L-170
L-230	Redwood	4000	Smith	L-230
L-260	Perryridge	1700	Hayes	L-155
loan			*borrower*	